

## ✓ Classify Waste Products Using Transfer Learning

### Table of contents

- [Classify Waste Products Using Transfer Learning](#)
  - [Introduction](#)
    - [Project Overview](#)
    - [Aim of the Project](#)
  - [Objectives](#)
    - [Tasks List](#)
    - [Sample Task: Sample screenshot showing code and output](#)
  - [Setup](#)
    - [Installing Required Libraries](#)
    - [Importing Required Libraries](#)
  - [Task 1: Print the version of tensorflow](#)
    - [Background](#)
    - [Create a model for distinguishing recyclable and organic waste images](#)
    - [Dataset](#)
    - [Importing Data](#)
    - [Define configuration options](#)
    - [Loading Images using ImageGeneratorClass](#)
      - [ImageDataGenerators](#)
  - [Task 2: Create a test generator using the test\\_datagen object](#)
  - [Task 3: Print the length of the train\\_generator](#)
    - [Pre-trained Models](#)
      - [VGG-16](#)
  - [Task 4: Print the summary of the model](#)
  - [Task 5: Compile the model](#)
    - [Fit and train the model](#)
    - [Plot loss curves for training and validation sets \(extract\\_feat\\_model\)](#)
  - [Task 6: Plot accuracy curves for training and validation sets \(extract\\_feat\\_model\)](#)
    - [Fine-Tuning model](#)
  - [Task 7: Plot loss curves for training and validation sets \(fine tune model\)](#)
  - [Task 8: Plot accuracy curves for training and validation sets \(fine tune model\)](#)
    - [Evaluate both models on test data](#)
  - [Task 9: Plot a test image using Extract Features Model \(index\\_to\\_plot = 1\)](#)
  - [Task 10: Plot a test image using Fine-Tuned Model \(index\\_to\\_plot = 1\)](#)
  - [Authors: Christopher Banner](#)

### [Overview](#)

EcoClean currently lacks an efficient and scalable method to automate the waste sorting process. The manual sorting of waste is not only labor-intensive but also prone to errors, leading to contamination of recyclable materials. The goal of this project is to leverage machine learning and computer vision to automate the classification of waste products, improving efficiency and reducing contamination rates. The project will use transfer learning with a pre-trained VGG16 model to classify images.

### [Aim of the Project](#)

The aim of the project is to develop an automated waste classification model that can accurately differentiate between recyclable and organic waste based on images. By the end of this project, you will have trained, fine-tuned, and evaluated a model using transfer learning, which can then be applied to real-world waste management processes.

**Final Output:** A trained model that classifies waste images into recyclable and organic categories.

## ✓ [Tasks List](#)

- Task 1: Print the version of tensorflow
- Task 2: Create a `test_generator` using the `test_datagen` object
- Task 3: Print the length of the `train_generator`
- Task 4: Print the summary of the model
- Task 5: Compile the model
- Task 6: Plot accuracy curves for training and validation sets (extract\_feat\_model)
- Task 7: Plot loss curves for training and validation sets (fine tune model)
- Task 8: Plot accuracy curves for training and validation sets (fine tune model)
- Task 9: Plot a test image using Extract Features Model (index\_to\_plot = 1)
- Task 10: Plot a test image using Fine-Tuned Model (index\_to\_plot = 1)

```
1 !pip install matplotlib
2 !pip install scikit-learn
3 !pip install tensorflow
4 !pip install keras
```

```
1 # Uninstall in a specific order to minimize temporary conflicts
2 !pip uninstall tensorflow -y
3 !pip uninstall tensorflow-gpu -y
4 !pip uninstall keras -y # Uninstall standalone Keras, if present
5 !pip uninstall h5py -y
6 !pip uninstall protobuf -y
7 !pip uninstall numpy -y # Uninstall numpy last of the major ones
8
9 print("Uninstall complete. Clearing pip cache...")
10 !pip cache purge
11 print("Installing compatible TensorFlow version (which will pull compatible numpy, h5py, protobuf)...")
12 !pip install tensorflow==2.16.1
```

## ▼ [Importing Required Libraries](#)

```
1 import numpy as np
2 import os
3 import random, shutil
4 import glob
5 from matplotlib import pyplot as plt
6 from matplotlib import pyplot
7 from matplotlib.image import imread
8 from os import makedirs, listdir
9 from shutil import copyfile
10 from random import seed
11 from random import random
12 os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
13 import os
14 import numpy as np
15 import matplotlib.pyplot as plt
16 import tensorflow as tf
17 from tensorflow import keras
18 from tensorflow.keras.preprocessing.image import ImageDataGenerator # This import should now work
19 from tensorflow.keras.applications import VGG16
20 from tensorflow.keras.models import Sequential
21 from tensorflow.keras.layers import Dense, Flatten, Dropout, BatchNormalization
22 from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint
23 from tensorflow.keras import optimizers
24 from tensorflow.keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D, Input
25 from tensorflow.keras.layers import Dense, Dropout, Flatten
26 from tensorflow.keras.models import Sequential, Model
27 from tensorflow.keras.applications import InceptionV3
28 from sklearn import metrics
29 import warnings
30 warnings.filterwarnings('ignore')
```

## ▼ [Print the version of tensorflow](#)

```
1 tf.__version__
```

```
↻ '2.16.1'
```

## ▼ [Background](#)

**Transfer learning** uses the concept of keeping the early layers of a pre-trained network, and re-training the later layers on a specific dataset. You can leverage some state of that network on a related task.

A typical transfer learning workflow in Keras looks something like this:

1. Initialize base model, and load pre-trained weights (e.g. ImageNet)
2. "Freeze" layers in the base model by setting `training = False`

3. Define a new model that goes on top of the output of the base model's layers.
4. Train resulting model on your data set.

## Create a model for distinguishing recyclable and organic waste images

### Dataset

[Waste Classification Dataset](#).

### Importing Data

This will create a `o-vs-r-split` directory in your environment.

```
1 import requests
2 import zipfile
3 from tqdm import tqdm
4
5 url = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/kd6057VPpABQ2FqCbg9VQ/o-vs-r-split-reduced-1200.zip"
6 file_name = "o-vs-r-split-reduced-1200.zip"
7
8 print("Downloading file")
9 with requests.get(url, stream=True) as response:
10     response.raise_for_status()
11     with open(file_name, 'wb') as f:
12         for chunk in response.iter_content(chunk_size=8192):
13             f.write(chunk)
14
15
16 def extract_file_with_progress(file_name):
17     print("Extracting file with progress")
18     with zipfile.ZipFile(file_name, 'r') as zip_ref:
19         members = zip_ref.infolist()
20         with tqdm(total=len(members), unit='file') as progress_bar:
21             for member in members:
22                 zip_ref.extract(member)
23                 progress_bar.update(1)
24     print("Finished extracting file")
25
26
27 extract_file_with_progress(file_name)
28
29 print("Finished extracting file")
30 os.remove(file_name)
```



```
Downloading file
Extracting file with progress
100%|██████████| 1207/1207 [01:32<00:00, 12.99file/s]
Finished extracting file
Finished extracting file
```

### Define configuration options

It's time to define some model configuration options.

- **batch size** is set to 32.
- The **number of classes** is 2.
- You will use 20% of the data for **validation** purposes.
- You have two **labels** in your dataset: organic (O), recyclable (R).

```
1 img_rows, img_cols = 150, 150
2 batch_size = 32
3 n_epochs = 10
4 n_classes = 2
5 val_split = 0.2
6 verbosity = 1
7 path = 'o-vs-r-split/train/'
8 path_test = 'o-vs-r-split/test/'
9 input_shape = (img_rows, img_cols, 3)
10 labels = ['O', 'R']
11 seed = 42
```

### Loading Images using ImageGeneratorClass

Transfer learning works best when models are trained on smaller datasets.

#### [ImageDataGenerators](#)

creating ImageDataGenerators used for training, validation and testing.

Image data generators create batches of tensor image data with real-time data augmentation. The generators loop over the data in batches and are useful in feeding data to the training process.

```


1 # Create ImageDataGenerators for training and validation and testing
2 train_datagen = ImageDataGenerator(
3     validation_split = val_split,
4     rescale=1.0/255.0,
5     width_shift_range=0.1,
6     height_shift_range=0.1,
7     horizontal_flip=True
8 )
9
10 val_datagen = ImageDataGenerator(
11     validation_split = val_split,
12     rescale=1.0/255.0,
13 )
14
15 test_datagen = ImageDataGenerator(
16     rescale=1.0/255.0
17 )

```

```

1 train_generator = train_datagen.flow_from_directory(
2     directory = path,
3     seed = seed,
4     batch_size = batch_size,
5     class_mode='binary',
6     shuffle = True,
7     target_size=(img_rows, img_cols),
8     subset = 'training'
9 )


```

 Found 800 images belonging to 2 classes.

```

1 val_generator = val_datagen.flow_from_directory(
2     directory = path,
3     seed = seed,
4     batch_size = batch_size,
5     class_mode='binary',
6     shuffle = True,
7     target_size=(img_rows, img_cols),
8     subset = 'validation'
9 )

```

 Found 200 images belonging to 2 classes.

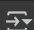
### ▼ Creating a `test_generator` using the `test_datagen` object

- **directory** should be set to `path_test`.
- **class\_mode** should be set to `'binary'`.
- **seed** should be set to `seed`.
- **batch\_size** should be set to `batch_size`.
- **shuffle** should be set to `False`.
- **target\_size** should be set to `(img_rows, img_cols)`.

```

1 # Task 2: Create a `test_generator` using the `test_datagen` object
2 test_generator = test_datagen.flow_from_directory(
3     directory = path,
4     seed = seed,
5     batch_size = batch_size,
6     class_mode='binary',
7     shuffle = True,
8     target_size=(img_rows, img_cols)
9 )

```


 Found 1000 images belonging to 2 classes.

### ▼ Print the length of the `train_generator`

```

1 # Task 3: print the length of the `train_generator`
2 test_generator = test_datagen.flow_from_directory(
3     directory = path,
4     seed = seed,
5     batch_size = batch_size,
6     class_mode='binary',
7     shuffle = True,
8     target_size=(img_rows, img_cols)
9 )

```

 Found 1000 images belonging to 2 classes.


Let's look at a few augmented images:

```

1 from pathlib import Path
2
3 IMG_DIM = (100, 100)
4
5 train_files = glob.glob('./o-vs-r-split/train/O/*')
6 train_files = train_files[:20]
7 train_imgs = [tf.keras.preprocessing.image.img_to_array(tf.keras.preprocessing.image.load_img(img, target_size=IMG_DIM)) for img in train_files]
8 train_imgs = np.array(train_imgs)
9 train_labels = [Path(fn).parent.name for fn in train_files]
10
11 img_id = 0
12 O_generator = train_datagen.flow(train_imgs[img_id:img_id+1], train_labels[img_id:img_id+1],
13                                 batch_size=1)
14 O = [next(O_generator) for i in range(0,5)]
15 fig, ax = plt.subplots(1,5, figsize=(16, 6))
16 print('Labels:', [item[1][0] for item in O])
17 l = [ax[i].imshow(O[i][0][0]) for i in range(0,5)]
18

```

Labels: ['O', 'O', 'O', 'O', 'O']



## Pre-trained Models

Pre-trained models are saved networks that have previously been trained on some large datasets. They are typically used for large-scale image-classification task. They can be used as they are or could be customized to a given task using transfer learning. These pre-trained models form the basis of transfer learning.

### VGG-16

Let us load the VGG16 model.

```

1 from tensorflow.keras.applications import vgg16
2
3 input_shape = (150, 150, 3)
4 vgg = vgg16.VGG16(include_top=False,
5                   weights='imagenet',
6                   input_shape=input_shape)
7
8

```

We flatten the output of a vgg model and assign it to the model `output`, we then use a Model object `basemodel` to group the layers into an object for training and inference . With the following inputs and outputs

inputs: `vgg.input`

outputs: `tf.keras.layers.Flatten()(output)`

```

1 output = vgg.layers[-1].output
2 output = tf.keras.layers.Flatten()(output)
3 basemodel = Model(vgg.input, output)

```

Next, you freeze the basemodel.

```

1 for layer in basemodel.layers:
2     layer.trainable = False

```

Create a new model on top. You add a Dropout layer for regularization, only these layers will change as for the lower layers you set `training=False` when calling the base model.

```

1 input_shape = basemodel.output_shape[1]
2
3 model = Sequential()
4 model.add(basemodel)
5 model.add(Dense(512, activation='relu'))
6 model.add(Dropout(0.3))

```

```

7 model.add(Dense(512, activation='relu'))
8 model.add(Dropout(0.3))
9 model.add(Dense(1, activation='sigmoid'))

```

## ▼ Print the summary of the model

```

1 # Task: print the summary of the model
2 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
functional (Functional)	(None, 8192)	14,714,688
dense (Dense)	(None, 512)	4,194,816
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262,656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 1)	513

Total params: 19,172,673 (73.14 MB)  
 Trainable params: 4,457,985 (17.01 MB)  
 Non-trainable params: 14,714,688 (56.13 MB)

## ▼ Compile the model

- **loss:** 'binary\_crossentropy'.
- **optimizer:** optimizers.RMSprop(learning\_rate=1e-4).
- **metrics:** ['accuracy'].

```

1 for layer in basemodel.layers:
2     layer.trainable = False
3
4 # Task 5: Compile the model
5 model.compile(
6     loss='binary_crossentropy',
7     optimizer=optimizers.RMSprop(learning_rate=1e-4),
8     metrics=['accuracy']
9 )

```

You will use early stopping to avoid over-training the model.

```

1 from tensorflow.keras.callbacks import LearningRateScheduler
2
3
4 checkpoint_path='O_R_tlearn_vgg16.keras'
5
6 # define step decay function
7 class LossHistory(tf.keras.callbacks.Callback):
8     def on_train_begin(self, logs={}):
9         self.losses = []
10        self.lr = []
11
12    def on_epoch_end(self, epoch, logs={}):
13        self.losses.append(logs.get('loss'))
14        self.lr.append(exp_decay(epoch))
15        print('lr:', exp_decay(len(self.losses)))
16
17 def exp_decay(epoch):
18     initial_lr = 1e-4
19     k = 0.1
20     lr = initial_lr * np.exp(-k*epoch)
21     return lr
22
23 # learning schedule callback
24 loss_history_ = LossHistory_()
25 lr_scheduler_ = LearningRateScheduler(exp_decay)
26
27 keras_callbacks = [
28     EarlyStopping(monitor = 'val_loss',
29                 patience = 4,
30                 mode = 'min',
31                 min_delta=0.01),
32     ModelCheckpoint(checkpoint_path, monitor='val_loss', save_best_only=True, mode='min')
33 ]
34
35 callbacks_list_ = [loss_history_, lr_scheduler_] + keras_callbacks

```

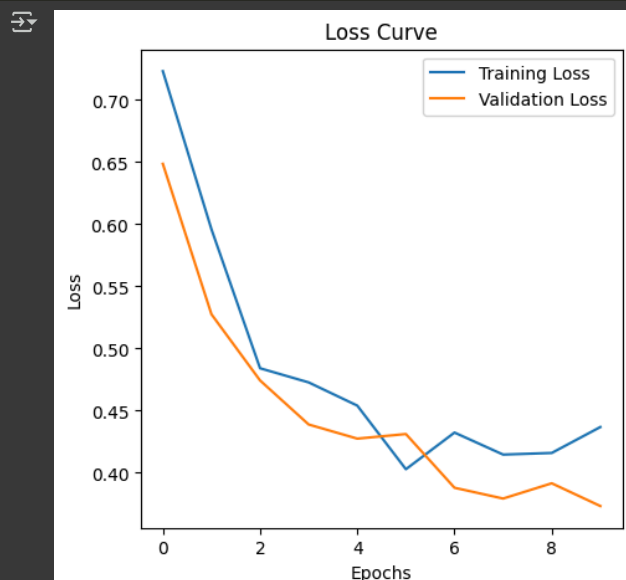
## Fit and train the model

```
1 extract_feat_model = model.fit(train_generator,  
2                               steps_per_epoch=5,  
3                               epochs=10,  
4                               callbacks = callbacks_list_,  
5                               validation_data=val_generator,  
6                               validation_steps=val_generator.samples // batch_size,  
7                               verbose=1)
```

```
Epoch 1/10  
lr: 9.048374180359596e-05 ----- 0s 5s/step - accuracy: 0.4627 - loss: 0.7538  
5/5 ----- 56s 12s/step - accuracy: 0.4700 - loss: 0.7487 - val_accuracy: 0.5729 - val_loss: 0.6486 - learning_rate: 1.0000e-04  
Epoch 2/10  
lr: 8.187307530779819e-05 ----- 0s 5s/step - accuracy: 0.6384 - loss: 0.6155  
5/5 ----- 52s 12s/step - accuracy: 0.6435 - loss: 0.6123 - val_accuracy: 0.7760 - val_loss: 0.5276 - learning_rate: 9.0484e-05  
Epoch 3/10  
lr: 7.408182206817179e-05 ----- 0s 5s/step - accuracy: 0.8142 - loss: 0.4771  
5/5 ----- 53s 12s/step - accuracy: 0.8118 - loss: 0.4783 - val_accuracy: 0.7812 - val_loss: 0.4743 - learning_rate: 8.1873e-05  
Epoch 4/10  
lr: 6.703200460356393e-05 ----- 0s 5s/step - accuracy: 0.7915 - loss: 0.4520  
5/5 ----- 52s 12s/step - accuracy: 0.7835 - loss: 0.4554 - val_accuracy: 0.8229 - val_loss: 0.4389 - learning_rate: 7.4082e-05  
Epoch 5/10  
lr: 6.065306597126335e-05 ----- 0s 5s/step - accuracy: 0.7755 - loss: 0.4868  
5/5 ----- 52s 12s/step - accuracy: 0.7817 - loss: 0.4814 - val_accuracy: 0.7969 - val_loss: 0.4275 - learning_rate: 6.7032e-05  
Epoch 6/10  
lr: 5.488116360940264e-05 ----- 0s 5s/step - accuracy: 0.8421 - loss: 0.3920  
5/5 ----- 51s 11s/step - accuracy: 0.8392 - loss: 0.3938 - val_accuracy: 0.8229 - val_loss: 0.4312 - learning_rate: 6.0653e-05  
Epoch 7/10  
lr: 4.9658530379140954e-05 ----- 0s 5s/step - accuracy: 0.7651 - loss: 0.4577  
5/5 ----- 51s 12s/step - accuracy: 0.7678 - loss: 0.4535 - val_accuracy: 0.8229 - val_loss: 0.3880 - learning_rate: 5.4881e-05  
Epoch 8/10  
lr: 4.493289641172216e-05 ----- 0s 5s/step - accuracy: 0.8369 - loss: 0.4091  
5/5 ----- 51s 12s/step - accuracy: 0.8349 - loss: 0.4100 - val_accuracy: 0.8281 - val_loss: 0.3793 - learning_rate: 4.9659e-05  
Epoch 9/10  
lr: 4.0656965974059915e-05 ----- 0s 5s/step - accuracy: 0.8465 - loss: 0.3891  
5/5 ----- 50s 11s/step - accuracy: 0.8439 - loss: 0.3936 - val_accuracy: 0.8073 - val_loss: 0.3916 - learning_rate: 4.4933e-05  
Epoch 10/10  
lr: 3.678794411714424e-05 ----- 0s 5s/step - accuracy: 0.8058 - loss: 0.4248  
5/5 ----- 51s 12s/step - accuracy: 0.8049 - loss: 0.4268 - val_accuracy: 0.8438 - val_loss: 0.3733 - learning_rate: 4.0657e-05
```

## Plot loss curves for training and validation sets (extract\_feat\_model)

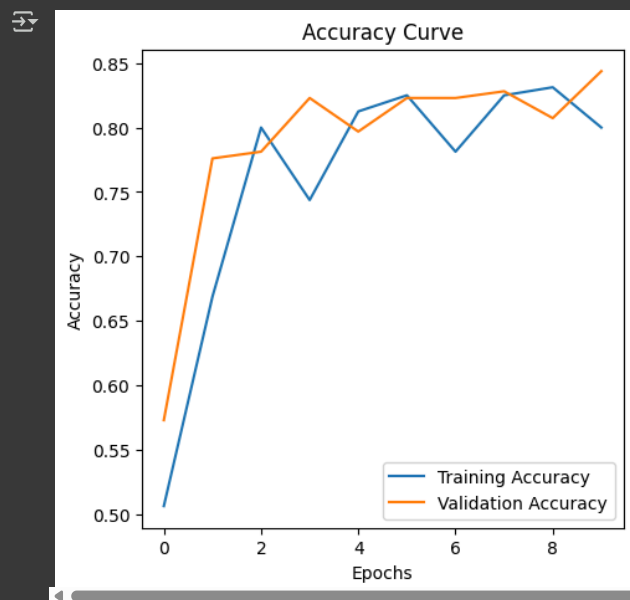
```
1 import matplotlib.pyplot as plt  
2  
3 history = extract_feat_model  
4  
5 # plot loss curve  
6 plt.figure(figsize=(5, 5))  
7 plt.plot(history.history['loss'], label='Training Loss')  
8 plt.plot(history.history['val_loss'], label='Validation Loss')  
9 plt.title('Loss Curve')  
10 plt.xlabel('Epochs')  
11 plt.ylabel('Loss')  
12 plt.legend()  
13  
14 plt.show()
```



## ✓ \*\* Plot accuracy curves for training and validation sets (extract\_feat\_model)\*\*

**NOTE:** As training is a stochastic process, the loss and accuracy graphs may differ across runs. As long as the general trend shows decreasing loss and increasing accuracy, the model is performing as expected and full marks will be awarded for the task.

```
1 import matplotlib.pyplot as plt
2
3 history = extract_feat_model
4 ## Task 6: Plot accuracy curves for training and validation sets
5 plt.figure(figsize=(5, 5))
6 plt.plot(history.history['accuracy'], label='Training Accuracy')
7 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
8 plt.title('Accuracy Curve')
9 plt.xlabel('Epochs')
10 plt.ylabel('Accuracy')
11 plt.legend()
12
13 plt.show()
```



## ✓ Fine-Tuning model

Fine-tuning is an optional step in transfer learning, it usually ends up improving the performance of the model.

You will **unfreeze** one layer from the base model and train the model again.

```
1 from tensorflow.keras.applications import vgg16
2
3 input_shape = (150, 150, 3)
4 vgg = vgg16.VGG16(include_top=False,
5                   weights='imagenet',
6                   input_shape=input_shape)
7 output = vgg.layers[-1].output
8 output = tf.keras.layers.Flatten()(output)
9 basemodel = Model(vgg.input, output)
10 for layer in basemodel.layers:
11     layer.trainable = False
12 display([layer.name for layer in basemodel.layers])
13 set_trainable = False
14 for layer in basemodel.layers:
15     if layer.name in ['block5_conv3']:
16         set_trainable = True
17     if set_trainable:
18         layer.trainable = True
19     else:
20         layer.trainable = False
21 for layer in basemodel.layers:
22     print(f"{layer.name}: {layer.trainable}")
```



```

☐ ['input_layer_2',
'block1_conv1',
'block1_conv2',
'block1_pool',
'block2_conv1',
'block2_conv2',
'block2_pool',
'block3_conv1',
'block3_conv2',
'block3_conv3',
'block3_pool',
'block4_conv1',
'block4_conv2',
'block4_conv3',
'block4_pool',
'block5_conv1',
'block5_conv2',
'block5_conv3',
'block5_pool',
'flatten_1']
input_layer_2: False
block1_conv1: False
block1_conv2: False
block1_pool: False
block2_conv1: False
block2_conv2: False
block2_pool: False
block3_conv1: False
block3_conv2: False
block3_conv3: False
block3_pool: False
block4_conv1: False
block4_conv2: False
block4_conv3: False
block4_pool: False
block5_conv1: False
block5_conv2: False
block5_conv3: True
block5_pool: True
flatten_1: True

```

Similar to what you did before, you will create a new model on top, and add a Dropout layer for regularization.

```

1 model = Sequential()
2 model.add(basemodel)
3 model.add(Dense(512, activation='relu'))
4 model.add(Dropout(0.3))
5 model.add(Dense(512, activation='relu'))
6 model.add(Dropout(0.3))
7 model.add(Dense(1, activation='sigmoid'))
8
9 checkpoint_path='0_R_tlearn_fine_tune_vgg16.keras'
10
11 # learning schedule callback
12 loss_history_ = LossHistory()
13 lr_rate_ = LearningRateScheduler(exp_decay)
14 keras_callbacks = [
15     EarlyStopping(monitor = 'val_loss',
16                   patience = 4,
17                   mode = 'min',
18                   min_delta=0.01),
19     ModelCheckpoint(checkpoint_path, monitor='val_loss', save_best_only=True, mode='min')
20 ]
21 callbacks_list_ = [loss_history_, lr_rate_] + keras_callbacks
22 model.compile(loss='binary_crossentropy',
23              optimizer=optimizers.RMSprop(learning_rate=1e-4),
24              metrics=['accuracy'])
25 fine_tune_model = model.fit(train_generator,
26                             steps_per_epoch=5,
27                             epochs=10,
28                             callbacks = callbacks_list_,
29                             validation_data=val_generator,
30                             validation_steps=val_generator.samples // batch_size,
31                             verbose=1)

```

```

☐ Epoch 1/10
lr: 9.048374180359596e-05----- 0s 5s/step - accuracy: 0.6036 - loss: 0.7620
5/5 ----- 60s 13s/step - accuracy: 0.6020 - loss: 0.7601 - val_accuracy: 0.6615 - val_loss: 0.5868 - learning_rate: 1.0000e-04
Epoch 2/10
lr: 8.187307530779819e-05----- 0s 5s/step - accuracy: 0.6868 - loss: 0.5760
5/5 ----- 56s 13s/step - accuracy: 0.6859 - loss: 0.5714 - val_accuracy: 0.8490 - val_loss: 0.4493 - learning_rate: 9.0484e-05
Epoch 3/10
lr: 7.408182206817179e-05----- 0s 5s/step - accuracy: 0.7725 - loss: 0.4945
5/5 ----- 53s 12s/step - accuracy: 0.7667 - loss: 0.4961 - val_accuracy: 0.7240 - val_loss: 0.5094 - learning_rate: 8.1873e-05
Epoch 4/10
lr: 6.703200460356393e-05----- 0s 5s/step - accuracy: 0.8190 - loss: 0.4782

```

```

5/5 ----- 57s 13s/step - accuracy: 0.8241 - loss: 0.4672 - val_accuracy: 0.8490 - val_loss: 0.3626 - learning_rate: 7.4082e-05
Epoch 5/10
lr: 6.065306597126335e-05 ----- 0s 5s/step - accuracy: 0.8631 - loss: 0.3537
5/5 ----- 53s 12s/step - accuracy: 0.8578 - loss: 0.3578 - val_accuracy: 0.7969 - val_loss: 0.3862 - learning_rate: 6.7032e-05
Epoch 6/10
lr: 5.488116360940264e-05 ----- 0s 5s/step - accuracy: 0.7985 - loss: 0.4742
5/5 ----- 56s 13s/step - accuracy: 0.8092 - loss: 0.4565 - val_accuracy: 0.8542 - val_loss: 0.3126 - learning_rate: 6.0653e-05
Epoch 7/10
lr: 4.9658530379140954e-05 ----- 0s 5s/step - accuracy: 0.8462 - loss: 0.3587
5/5 ----- 54s 12s/step - accuracy: 0.8479 - loss: 0.3574 - val_accuracy: 0.8750 - val_loss: 0.3071 - learning_rate: 5.4881e-05
Epoch 8/10
lr: 4.493289641172216e-05 ----- 0s 5s/step - accuracy: 0.8672 - loss: 0.3206
5/5 ----- 54s 12s/step - accuracy: 0.8633 - loss: 0.3289 - val_accuracy: 0.8802 - val_loss: 0.2994 - learning_rate: 4.9659e-05
Epoch 9/10
lr: 4.0656965974059915e-05 ----- 0s 5s/step - accuracy: 0.8066 - loss: 0.3703
5/5 ----- 55s 12s/step - accuracy: 0.8065 - loss: 0.3782 - val_accuracy: 0.8750 - val_loss: 0.2913 - learning_rate: 4.4933e-05
Epoch 10/10
lr: 3.678794411714424e-05 ----- 0s 5s/step - accuracy: 0.8839 - loss: 0.2748
5/5 ----- 55s 13s/step - accuracy: 0.8824 - loss: 0.2766 - val_accuracy: 0.8646 - val_loss: 0.2849 - learning_rate: 4.0657e-05

```

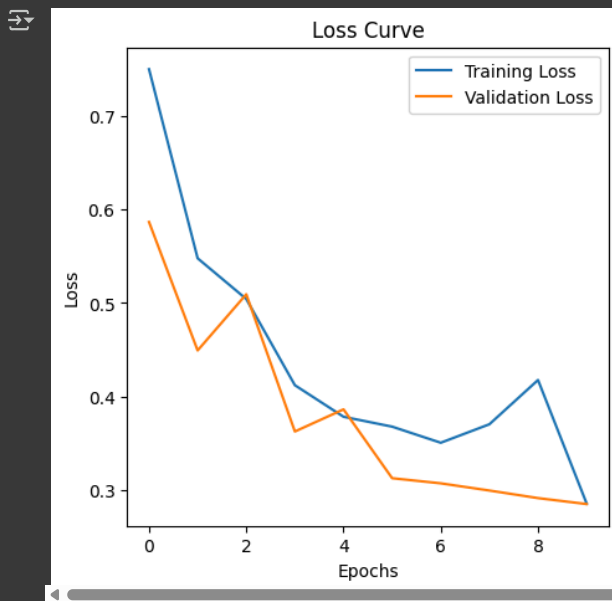
## ✚ \*\* Plot loss curves for training and validation sets (fine tune model)\*\*

**NOTE:** As training is a stochastic process, the loss and accuracy graphs may differ across runs. As long as the general trend shows decreasing loss and increasing accuracy, the model is performing as expected and full marks will be awarded for the task.

```

1 import matplotlib.pyplot as plt
2
3 history = fine_tune_model
4
5 ## Task 7: Plot loss curves for training and validation sets (fine tune model)
6
7 plt.figure(figsize=(5, 5))
8 plt.plot(history.history['loss'], label='Training Loss')
9 plt.plot(history.history['val_loss'], label='Validation Loss')
10 plt.title('Loss Curve')
11 plt.xlabel('Epochs')
12 plt.ylabel('Loss')
13 plt.legend()
14
15 plt.show()

```



## ✚ \*\* Plot accuracy curves for training and validation sets (fine tune model)\*\*

**NOTE:** As training is a stochastic process, the loss and accuracy graphs may differ across runs. As long as the general trend shows decreasing loss and increasing accuracy, the model is performing as expected and full marks will be awarded for the task.

```

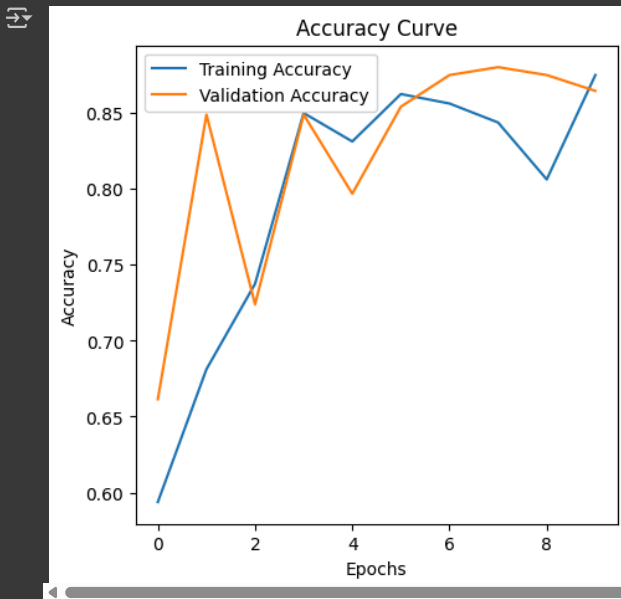
1 import matplotlib.pyplot as plt
2 history = fine_tune_model
3
4 # Task 8: Plot accuracy curves for training and validation sets (fine tune model)
5 plt.figure(figsize=(5, 5))
6 plt.plot(history.history['accuracy'], label='Training Accuracy')
7 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
8 plt.title('Accuracy Curve')
9 plt.xlabel('Epochs')
10 plt.ylabel('Accuracy')

```

```

11 plt.legend()
12
13 plt.show()

```



## Evaluate both models on test data

- Load saved models
- Load test images
- Make predictions for both models
- Convert predictions to class labels
- Print classification report for both models

```

1 from pathlib import Path
2
3 # Load saved models
4 extract_feat_model = tf.keras.models.load_model('O_R_tlearn_vgg16.keras')
5 fine_tune_model = tf.keras.models.load_model('O_R_tlearn_fine_tune_vgg16.keras')
6
7 IMG_DIM = (150, 150)
8
9 # Load test images
10 test_files_O = glob.glob('./o-vs-r-split/test/O/*')
11 test_files_R = glob.glob('./o-vs-r-split/test/R/*')
12 test_files = test_files_O[:50] + test_files_R[:50]
13
14 test_imgs = [tf.keras.preprocessing.image.img_to_array(tf.keras.preprocessing.image.load_img(img, target_size=IMG_DIM)) for img in test_files]
15 test_imgs = np.array(test_imgs)
16 test_labels = [Path(fn).parent.name for fn in test_files]
17
18 # Standardize
19 test_imgs_scaled = test_imgs.astype('float32')
20 test_imgs_scaled /= 255
21
22 class2num_lt = lambda l: [0 if x == 'O' else 1 for x in l]
23 num2class_lt = lambda l: ['O' if x < 0.5 else 'R' for x in l]
24
25 test_labels_enc = class2num_lt(test_labels)
26
27 # Make predictions for both models
28 predictions_extract_feat_model = extract_feat_model.predict(test_imgs_scaled, verbose=0)
29 predictions_fine_tune_model = fine_tune_model.predict(test_imgs_scaled, verbose=0)
30
31 # Convert predictions to class labels
32 predictions_extract_feat_model = num2class_lt(predictions_extract_feat_model)
33 predictions_fine_tune_model = num2class_lt(predictions_fine_tune_model)
34
35 # Print classification report for both models
36 print('Extract Features Model')
37 print(metrics.classification_report(test_labels, predictions_extract_feat_model))
38 print('Fine-Tuned Model')
39 print(metrics.classification_report(test_labels, predictions_fine_tune_model))
40

```

```

Extract Features Model
      precision    recall  f1-score   support

O         0.70      0.90      0.79        50
R         0.86      0.62      0.72        50

```

accuracy			0.76	100
macro avg	0.78	0.76	0.76	100
weighted avg	0.78	0.76	0.76	100

Fine-Tuned Model				
	precision	recall	f1-score	support
O	0.78	0.90	0.83	50
R	0.88	0.74	0.80	50

accuracy			0.82	100
macro avg	0.83	0.82	0.82	100
weighted avg	0.83	0.82	0.82	100

```

1 # Plot one of the images with actual label and predicted label as title
2 def plot_image_with_title(image, model_name, actual_label, predicted_label):
3     plt.imshow(image)
4     plt.title(f"Model: {model_name}, Actual: {actual_label}, Predicted: {predicted_label}")
5     plt.axis('off')
6     plt.show()
7
8 # Specify index of image to plot, for example index 0
9 index_to_plot = 0
10 plot_image_with_title(
11     image=test_imgs[index_to_plot].astype('uint8'),
12     model_name='Extract Features Model',
13     actual_label=test_labels[index_to_plot],
14     predicted_label=predictions_extract_feat_model[index_to_plot],
15 )

```



Model: Extract Features Model, Actual: O, Predicted: O



#### ▼ **\*\* Plot a test image using Extract Features Model (index\_to\_plot = 1)\*\***

**NOTE:** Due to the inherent nature of neural networks, predictions may vary from the actual labels. For instance, if the actual label is 'O', the prediction could be either 'O' or 'R', both of which are possible outcomes, and full marks will be awarded for the task.

```

1 # Task 9: Plot a test image using Extract Features Model (index_to_plot = 1)
2
3 # Specify index of image to plot
4 index_to_plot = 1
5 plot_image_with_title(
6     image=test_imgs[index_to_plot].astype('uint8'),
7     model_name='Extract Features Model',
8     actual_label=test_labels[index_to_plot],
9     predicted_label=predictions_extract_feat_model[index_to_plot],
10 )

```



Model: Extract Features Model, Actual: O, Predicted: O



✓ Plot a test image using Fine-Tuned Model (index\_to\_plot = 1)

**NOTE:** Due to the inherent nature of neural networks, predictions may vary from the actual labels. For instance, if the actual label is 'O', the prediction could also be 'O' or 'R', both of which are possible outcomes, and full marks will be awarded for the task.

```
1 # Task 10: Plot a test image using Fine-Tuned Model (index_to_plot = 1)
2
3 # Specify index of image to plot
4 index_to_plot = 1
5 plot_image_with_title(
6     image=test_imgs[index_to_plot].astype('uint8'),
7     model_name='Fine-Tuned Model',
8     actual_label=test_labels[index_to_plot],
9     predicted_label=predictions_fine_tune_model[index_to_plot],
10 )
```



Model: Fine-Tuned Model, Actual: O, Predicted: O

